

```
import pandas as pd

df=pd.read_csv("Social_Network_Ads.csv")
df
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1
397	15654296	Female	50	20000	1
398	15755018	Male	36	33000	0
399	15594041	Female	49	36000	1

400 rows × 5 columns

```
# x=df.drop(columns='Purchased')
x=df.drop('Purchased',axis=1)
y=df['Purchased']
print("\n",x)
print("\n",y)
x=x.drop(columns='Gender')
print("\n",x)
```

	User ID	Gender	Age	EstimatedSalary
0	15624510	Male	19	19000
1	15810944	Male	35	20000
2	15668575	Female	26	43000
3	15603246	Female	27	57000
4	15804002	Male	19	76000
..
395	15691863	Female	46	41000
396	15706071	Male	51	23000
397	15654296	Female	50	20000
398	15755018	Male	36	33000
399	15594041	Female	49	36000

[400 rows x 4 columns]

0	0
1	0
2	0
3	0
4	0
..	
395	1
396	1
397	1
398	0
399	1

Name: Purchased, Length: 400, dtype: int64

	User ID	Age	EstimatedSalary
0	15624510	19	19000
1	15810944	35	20000
2	15668575	26	43000
3	15603246	27	57000
4	15804002	19	76000
..
395	15691863	46	41000
396	15706071	51	23000
397	15654296	50	20000
398	15755018	36	33000
399	15594041	49	36000

[400 rows x 3 columns]

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=1)
```

```
print(x_train)
```

```
print(x_test)
```

```
print(y_train)
```

```
print(y_test)
```



	User ID	Age	EstimatedSalary
39	15782806	27	31000
167	15614827	35	71000
383	15707634	49	28000
221	15663161	35	91000
351	15591279	37	75000
..
255	15750056	52	90000
72	15595228	20	23000
396	15706071	51	23000
235	15646227	46	79000
37	15689425	30	49000

[280 rows x 3 columns]

	User ID	Age	EstimatedSalary
398	15755018	36	33000
125	15697020	39	61000
328	15796351	36	118000
339	15665760	39	122000
172	15794661	26	118000
..
91	15636428	30	116000
322	15674331	41	52000
248	15730688	41	52000
186	15724402	20	82000
395	15691863	46	41000

[120 rows x 3 columns]

39	0
167	0
383	1
221	1
351	0
..	..
255	1
72	0
396	1
235	1
37	0

Name: Purchased, Length: 280, dtype: int64

```
398    0
125    0
328    1
339    1
172    0
..
91     0
322    0
248    0
186    0
395    1
Name: Purchased, Length: 120, dtype: int64
```

```
from sklearn.linear_model import LogisticRegression
```

```
model1=LogisticRegression()
model1.fit(x_train,y_train)
Pred_y=model1.predict(x_test)
print(Pred_y)
```

```
[0 0 1 1 1 0 0 0 1 0 0 0 1 0 1 0 0 0 1 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0
 0 0 0 1 1 1 0 0 1 0 1 0 0 0 0 0 0 0 0 0 1 1 0 0 1 0 1 1 0 0 1 0 1 0 1 0
 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0
 0 0 0 1 1 0 0 0 0]
```

```
from sklearn.metrics import accuracy_score
```

```
acc=accuracy_score(y_test,Pred_y)
print(acc)
```

```
0.6916666666666667
```

```
from sklearn.preprocessing import StandardScaler
```

```
std=StandardScaler()
x_test=std.fit_transform(x_test)
x_train=std.fit_transform(x_train)
print(x_test)
print(x_train)
```

```
[[ 1.06520773e+00 -2.75451436e-01 -1.14462515e+00]
 [ 2.16654071e-01  4.66866840e-03 -3.08732842e-01]
 [ 1.66994011e+00 -2.75451436e-01  1.39290506e+00]
 [-2.40702862e-01  4.66866840e-03  1.51231825e+00]
 [ 1.64521416e+00 -1.20918512e+00  1.39290506e+00]
 [ 5.17169471e-01 -8.87046996e-02 -1.89319656e-01]
 [-2.88058544e-02 -1.76942532e+00 -1.05506526e+00]
 [-1.25370604e+00  9.38402348e-01  5.27159462e-01]
 [-1.83701499e-01 -7.42318275e-01 -1.59242460e+00]
 [-6.37400748e-01  8.45028980e-01  2.07953089e+00]
 [-5.54737066e-01 -4.62198171e-01  1.96534203e-02]
 [-7.65580661e-01  4.66866840e-03  4.95067168e-02]
 [-1.17131632e-01 -3.68824803e-01  1.96534203e-02]
 [-1.56497503e+00  8.45028980e-01  1.78099792e+00]
 [-1.03908772e+00  1.31189582e+00  3.18186386e-01]
 [ 4.58548120e-02  1.59201592e+00  1.84070451e+00]
 [ 1.35381371e+00  1.96550940e+00  3.48039683e-01]
 [ 6.81223959e-01 -1.11581175e+00 -3.98292732e-01]
 [-2.18683599e-01 -1.02243838e+00  4.67452869e-01]
 [-1.29477453e+00  1.96550940e+00  9.15252318e-01]
 [ 6.17580240e-01  9.80420364e-02  1.09213310e-01]
 [ 1.43908166e+00  1.03177572e+00  4.97306166e-01]
```

```
[ 4.88815097e-01  4.71535508e-01  2.01982429e+00]
[-8.73950552e-01  7.51655612e-01 -8.46092181e-01]
[-9.47045730e-01  5.64908876e-01 -1.35359822e+00]
[ 1.86704544e+00 -1.20918512e+00 -1.68198449e+00]
[-1.57777693e+00  1.77876266e+00 -7.26678995e-01]
[ 3.80167222e-01  9.38402348e-01  7.93600134e-02]
[ 1.41359491e+00  1.31189582e+00 -1.11477185e+00]
[ 1.64382424e+00  1.21852245e+00  1.27349188e+00]
[-2.02297170e-01  4.66866840e-03 -8.75945477e-01]
[ 1.78187990e+00 -1.86279869e+00  1.39066607e-01]
[-1.51389912e+00 -1.95617206e+00  4.37599573e-01]
[ 6.62116213e-01  1.68538929e+00  7.93600134e-02]
[-1.23143805e+00 -1.11581175e+00  3.77892979e-01]
[-1.26255764e+00 -8.35691643e-01  2.58479793e-01]
[-5.75381041e-01 -1.58267859e+00 -1.59242460e+00]
[-1.42893841e+00 -6.48944907e-01  4.37599573e-01]
[-4.08429666e-01  1.03177572e+00 -1.53271800e+00]
[ 8.61489307e-01 -1.86279869e+00 -1.38345152e+00]
[ 9.36237758e-01  7.51655612e-01  2.16909077e+00]
[-3.12305704e-01  1.77876266e+00  8.85399021e-01]
[-1.62810668e+00 -4.62198171e-01  1.30334517e+00]
[-1.27399888e+00 -1.48930522e+00 -1.59466359e-01]
[ 1.36246048e+00  1.59201592e+00 -3.38586139e-01]
[-2.67345440e-01 -7.42318275e-01  1.39290506e+00]
[-2.00307389e-01  8.45028980e-01 -1.08491855e+00]
[ 7.57303808e-01  7.51655612e-01  1.24363858e+00]
[ 9.49317640e-01  4.66866840e-03  2.28626496e-01]
[-1.01241588e+00  1.21852245e+00 -9.95358664e-01]
[ 5.24645780e-01 -1.39593185e+00 -3.98292732e-01]
[-8.78442190e-01 -1.82078068e-01 -5.47559215e-01]
[ 5.67601632e-01  2.84788772e-01  2.58479793e-01]
[ 3.27306358e-01  6.58282244e-01 -1.29389163e+00]
[-1.26890336e-01  2.84788772e-01  4.95067168e-02]
[-7.48243234e-01 -1.82078068e-01 -2.78879546e-01]
[ 1.59710829e+00  1.96550940e+00 -8.75945477e-01]
[-1.32816188e+00 -2.75451436e-01 -5.77413512e-01]
```

```
model2=LogisticRegression()
model2.fit(x_train,y_train)
Pred_y=model2.predict(x_test)
print(Pred_y)
```

```
[0 0 1 1 0 0 0 1 0 1 0 0 0 1 1 1 1 0 0 1 0 1 1 0 0 0 1 1 1 1 0 0 0 1 0 0 0
 0 0 0 1 1 0 0 1 0 0 1 0 1 0 0 0 0 0 0 1 0 1 1 0 0 1 0 0 1 0 0 1 0 1 0 1 0
 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 1 1 0 0 0 1 1 0 0 0 0 1 0 0 0 0 0 1 1 0 0
 0 0 0 0 0 0 0 0 0]
```

```
acc2=accuracy_score(y_test,Pred_y)
print(acc2)
```

```
0.8
```


```
from sklearn.metrics import confusion_matrix
from sklearn.metrics import ConfusionMatrixDisplay
# in order to evaluate model when algorithm is classification based
```

```
cm=confusion_matrix(y_test,Pred_y)
cm
```

```
array([[65,  7],
       [17, 31]], dtype=int64)
```

```
ConfusionMatrixDisplay(cm).plot()
```

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x125323e3250>



```
from sklearn.metrics import recall_score, precision_score
```

```
print('recall = ', recall_score(y_test, Pred_y))
```

```
recall = 0.6458333333333334
```